

Notebook

August 1, 2018

Contents

1	Incorporating Textual Features	2
1.0.1	Problem	4
1.0.2	Tokenizing Text	4
1.0.3	Using the essay as features	7
1.0.4	min_df	8
1.0.5	Stop Words	9
1.0.6	tf-idf	10
1.0.7	Problem	11
1.0.8	n-Grams	11

List of Figures

List of Tables

List of Codes

1 Incorporating Textual Features

Continuing with the Donor's Choose example, we will examine how to make use of the textual information in columns like the `project_essay_1` column.

```
1 %matplotlib inline
2 import matplotlib.pyplot as plt
3 import pandas as pd
4 import numpy as np
5
6 from sklearn.model_selection import GridSearchCV
7 from sklearn.linear_model import LogisticRegression
8 from sklearn.pipeline import make_pipeline
9 from sklearn.metrics import classification_report
```

```
1 d_train = pd.read_csv('data/train.csv')
```

```
1 d_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 182080 entries, 0 to 182079
Data columns (total 16 columns):
id                                     182080 non-null object
teacher_id                           182080 non-null object
teacher_prefix                        182076 non-null object
school_state                          182080 non-null object
project_submitted_datetime           182080 non-null object
project_grade_category               182080 non-null object
project_subject_categories            182080 non-null object
project_subject_subcategories         182080 non-null object
project_title                        182080 non-null object
project_essay_1                      182080 non-null object
project_essay_2                      182080 non-null object
project_essay_3                      6374 non-null object
project_essay_4                      6374 non-null object
project_resource_summary              182080 non-null object
teacher_number_of_previously_posted_projects 182080 non-null int64
project_is_approved                  182080 non-null int64
dtypes: int64(2), object(14)
memory usage: 22.2+ MB
```

```
1 d_train.project_essay_1[0]
```

```
'Most of my kindergarten students come from low-income households and
are considered \"at-risk\". These kids walk to school alongside
their parents and most have never been further than walking distance
from their house. For 80% of my students, English is not their first
language or the language spoken at home.
While my
kindergarten kids have many obstacles in front of them, they come to
school each day excited and ready to learn. Most students started the
year out never being in a school setting. At the start of the year
many had never been exposed to letters. Each day they soak up more
knowledge and try their hardest to succeed. They are highly motivated
to learn new things every day. We are halfway through the year and
they are starting to take off. They know know all letters, some sight
```

words, numbers to 20, and a majority of their letter sounds because of their hard work and determination. I am excited to see the places we will go from here!'

```
1 essay_sample = d_train.project_essay_1[0]
```

```
1 for i in essay_sample[:10]:  
2     print(i)
```

M
o
s
t

o
f

m
y

```
1 sent = ["This is a      very boring      example      . ",  
2         "Now maybe another.  ",  
3         "  Will this be different ? "]
```

```
1 [string.strip() for string in sent]
```

```
['This is a      very boring      example      .',  
'Now maybe another.',  
'Will this be different ?']
```

```
1 [string.replace('.', ' ') for string in sent]
```

```
['This is a      very boring      example      ',  
'Now maybe another  ',  
'  Will this be different ? ']
```

```
1 [string.lower() for string in sent]
```

```
['this is a      very boring      example      . ',  
'now maybe another.  ',  
'  will this be different ? ']
```

```
1 import re
```

```
1 [re.sub(r"[aeiouAEIOU]", "X", string) for string in sent]
```

```
[ 'ThXs Xs X      vXry bXrXng      XxXmplX      . ',
  'NXw mXybX XnXthXr. ',
  '  WXll thXs bX dXffXrXnt ? ']
```

1.0.1 Problem

Using the `essay_sample` variable (first essay from first row of our Donor's Choose data), use the basic text strategies to do the following:

- remove any punctuation, if important to nature of word use (! vs. ?) determine a way to account for this.
- make sure all words are lowercase
- choose a few words that you believe to be the most important in the essay. Why did you choose these?

1.0.2 Tokenizing Text

```
1 from sklearn.feature_extraction.text import CountVectorizer
```

```
1 vect = CountVectorizer()
```

```
1 vect.fit(sent)
```

```
CountVectorizer(analyzer='word', binary=False, decode_error='strict',
                 dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
                 lowercase=True, max_df=1.0, max_features=None, min_df=1,
                 ngram_range=(1, 1), preprocessor=None, stop_words=None,
                 strip_accents=None, token_pattern='(?u)\\b\\w+\\b',
                 tokenizer=None, vocabulary=None)
```

```
1 vect.vocabulary_
```

```
{'another': 0,
 'be': 1,
 'boring': 2,
 'different': 3,
 'example': 4,
 'is': 5,
 'maybe': 6,
 'now': 7,
 'this': 8,
 'very': 9,
 'will': 10}
```

```
1 vect.transform(sent)
```

```
<3x11 sparse matrix of type '<class 'numpy.int64'>'
  with 12 stored elements in Compressed Sparse Row format>
```

```
1 bag = vect.transform(sent)
2 print("Content {}".format(repr(bag)))
```

```
Content <3x11 sparse matrix of type '<class 'numpy.int64'>'
with 12 stored elements in Compressed Sparse Row format>
```

```
1 bag.toarray()
```

```
array([[0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0],
       [1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0],
       [0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1]])
```

```
1 vect.transform(sent)
```

```
<3x11 sparse matrix of type '<class 'numpy.int64'>'
with 12 stored elements in Compressed Sparse Row format>
```

```
1 vect.vocabulary_
```

```
{'another': 0,
 'be': 1,
 'boring': 2,
 'different': 3,
 'example': 4,
 'is': 5,
 'maybe': 6,
 'now': 7,
 'this': 8,
 'very': 9,
 'will': 10}
```

```
1 sent
```

```
['This is a      very boring      example      . ',
 'Now maybe another.  ',
 ' Will this be different ? ']
```

```
1 example_slice = d_train[:50]
2 example_slice.head()
```

```
      id      teacher_id teacher_prefix
school_state \
0  p036502  484aaf11257089a66cfedc9461c6bd0a      Ms.
NV
1  p039565  df72a3ba8089423fa8a94be88060f6ed      Mrs.
GA
2  p233823  a9b876a9252e08a55e3d894150f75ba3      Ms.
UT
3  p185307  525fdbb6ec7f538a48beebaa0a51b24f      Mr.
NC
4  p013780  a63b5547a7239eae4c1872670848e61a      Mr.
CA

project_submitted_datetime project_grade_category \
```

```

0      2016-11-18 14:45:59      Grades PreK-2
1      2017-04-26 15:57:28      Grades 3-5
2      2017-01-01 22:57:44      Grades 3-5
3      2016-08-12 15:42:11      Grades 3-5
4      2016-08-06 09:09:11      Grades 6-8

      project_subject_categories \
0          Literacy & Language
1      Music & The Arts, Health & Sports
2      Math & Science, Literacy & Language
3          Health & Sports
4          Health & Sports

      project_subject_subcategories \
0          Literacy
1          Performing Arts, Team Sports
2      Applied Sciences, Literature & Writing
3          Health & Wellness
4          Health & Wellness

      project_title \
0          Super Sight Word Centers
1          Keep Calm and Dance On
2          Lets 3Doodler to Learn
3      \"Kid Inspired\" Equipment to Increase Activit...
4      We need clean water for our culinary arts class!

      project_essay_1 \
0      Most of my kindergarten students come from low...
1      Our elementary school is a culturally rich sch...
2      Hello;\r\nMy name is Mrs. Brotherton. I teach ...
3      My students are the greatest students but are ...
4      My students are athletes and students who are ...

      project_essay_2 project_essay_3
\
0      I currently have a differentiated sight word c...      NaN
1      We strive to provide our diverse population of...      NaN
2      We are looking to add some 3Doodler to our cla...      NaN
3      The student's project which is totally \"kid-i...      NaN
4      For some reason in our kitchen the water comes...      NaN

project_essay_4      project_resource_summary
\
0      NaN      My students need 6 Ipod Nano's to create and d...
1      NaN      My students need matching shirts to wear for d...
2      NaN      My students need the 3doodler. We are an SEM s...
3      NaN      My students need balls and other activity equi...
4      NaN      My students need a water filtration system for...

      teacher_number_of_previously_posted_projects      project_is_approved
0          26          1
1          1          0
2          5          1
3          16          0

```

```
1 vect.fit(example_slice.project_essay_1)
```

```
CountVectorizer(analyzer='word', binary=False, decode_error='strict',
               dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
               lowercase=True, max_df=1.0, max_features=None, min_df=5,
               ngram_range=(1, 1), preprocessor=None, stop_words='english',
               strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',
               tokenizer=None, vocabulary=None)
```

```
1 X = vect.transform(example_slice.project_essay_1)
```

```
1 feature_names = vect.get_feature_names()
```

```
1 feature_names[:10]
```

```
['able',
 'academically',
 'achieve',
 'activities',
 'area',
 'autism',
 'backgrounds',
 'best',
 'breakfast',
 'challenges']
```

```
1 feature_names[-10:]
```

```
['title',
 'use',
 'variety',
 'want',
 'way',
 'work',
 'working',
 'world',
 'year',
 'years']
```

```
1 len(feature_names)
```

```
103
```

1.0.3 Using the essay as features

```
1 from sklearn.linear_model import LogisticRegression
```

```
1 X_train = X[:40]
2 X_test = X[-10:]
3 y_train = example_slice.project_is_approved[:40]
4 y_test = example_slice.project_is_approved[-10:]
```

```
1 clf = LogisticRegression()
2 clf.fit(X_train, y_train)
3 clf.score(X_test, y_test)
```

0.7

```
1 clf = LogisticRegression()
2 params = {'C': [0.1, 1.0, 5.0, 10, 100, 1000]}
3 grid = GridSearchCV(clf, param_grid = params)
4 grid.fit(X_train, y_train)
```

```
GridSearchCV(cv=None, error_score='raise',
             estimator=LogisticRegression(C=1.0, class_weight=None, dual=
False, fit_intercept=True,
             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs
=1,
             penalty='l2', random_state=None, solver='liblinear', tol
=0.0001,
             verbose=0, warm_start=False),
             fit_params=None, iid=True, n_jobs=1,
             param_grid={'C': [0.1, 1.0, 5.0, 10, 100, 1000]}},
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring=None, verbose=0)
```

```
1 grid.best_estimator_.score(X_test, y_test)
```

0.8

```
1 grid.best_params_
```

```
{'C': 0.1}
```

1.0.4 min_df

When building the vocabulary ignore terms that have a document frequency strictly lower than the given threshold. This value is also called cut-off in the literature. If float, the parameter represents a proportion of documents, integer absolute counts. This parameter is ignored if vocabulary is not None.

```
1 vect = CountVectorizer(min_df=3)
2 vect.fit(example_slice.project_essay_1)
3 X = vect.transform(example_slice.project_essay_1)
```



```

1 X_train = X[:40]
2 X_test = X[-10:]
3 y_train = example_slice.project_is_approved[:40]
4 y_test = example_slice.project_is_approved[-10:]

```

```

1 clf = LogisticRegression()
2 params = {'C': [0.1, 1.0, 5.0, 10, 100, 1000]}
3 grid = GridSearchCV(clf, param_grid = params)
4 grid.fit(X_train, y_train)

```

```

GridSearchCV(cv=None, error_score='raise',
             estimator=LogisticRegression(C=1.0, class_weight=None, dual=
False, fit_intercept=True,
             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs
=1,
             penalty='l2', random_state=None, solver='liblinear', tol
=0.0001,
             verbose=0, warm_start=False),
             fit_params=None, iid=True, n_jobs=1,
             param_grid={'C': [0.1, 1.0, 5.0, 10, 100, 1000]}},
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring=None, verbose=0)

```

```

1 grid.best_estimator_.score(X_test, y_test)

```

0.8

1.0.5 Stop Words

```

1 from sklearn.feature_extraction.text import ENGLISH_STOP_WORDS

```

```

1 list(ENGLISH_STOP_WORDS)[:10]

```

```

['nothing',
 'are',
 'they',
 'above',
 'thru',
 'forty',
 'other',
 'eight',
 'ie',
 'third']

```

```

1 vect = CountVectorizer(min_df=5, stop_words='english')

```

```

1 vect.fit(example_slice.project_essay_1)
2 X = vect.transform(example_slice.project_essay_1)

```

```

1 X_train = X[:40]
2 X_test = X[-10:]
3 y_train = example_slice.project_is_approved[:40]
4 y_test = example_slice.project_is_approved[-10:]

```

```

1 clf = LogisticRegression()
2 params = {'C': [0.1, 1.0, 5.0, 10, 100, 1000]}
3 grid = GridSearchCV(clf, param_grid = params)
4 grid.fit(X_train, y_train)

```

```

GridSearchCV(cv=None, error_score='raise',
             estimator=LogisticRegression(C=1.0, class_weight=None, dual=
False, fit_intercept=True,
             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs
=1,
             penalty='l2', random_state=None, solver='liblinear', tol
=0.0001,
             verbose=0, warm_start=False),
             fit_params=None, iid=True, n_jobs=1,
             param_grid={'C': [0.1, 1.0, 5.0, 10, 100, 1000]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring=None, verbose=0)

```

```

1 grid.best_estimator_.score(X_test, y_test)

```

0.8

1.0.6 tf-idf

```

1 from sklearn.feature_extraction.text import TfidfVectorizer

```

```

1 pipe = make_pipeline(TfidfVectorizer(), LogisticRegression())

```

```

1 tfi = TfidfVectorizer()
2 X = example_slice.project_essay_1
3 tfi.fit(X)
4 X_tfidf = tfi.transform(X)

```

```

1 X_tfidf.toarray()[:10]

```

```

array([[0.         , 0.         , 0.         , ..., 0.         , 0.
,
       0.         ],
      [0.         , 0.         , 0.         , ..., 0.         , 0.
,
       0.         ],
      [0.         , 0.         , 0.         , ..., 0.         , 0.
,
       0.17757928],
      ...,
      [0.         , 0.         , 0.         , ..., 0.25583394, 0.
,
       0.         ],

```

```

        [0.          , 0.          , 0.          , ..., 0.          , 0.
,
        0.          ],
        [0.          , 0.          , 0.          , ..., 0.          , 0.
,
        0.          ]])

```

```

1 max_vals = X_tfidf.max(axis=0).toarray().ravel()

```

```

1 feature_names = np.array(tfi.get_feature_names())

```

```

1 tfidf_sorted = max_vals.argsort()

```

```

1 feature_names[tfidf_sorted[:10]]

```

```

array(['00', 'american', 'since', 'shelving', 'round', 'asking', '
proper',
       'nthis', 'monday', 'look'], dtype='<U17')

```

```

1 feature_names[tfidf_sorted[-10:]]

```

```

array(['highly', 'physical', 'to', 'world', 'band', 'computer', 'town
',
       'very', 'me', 'delays'], dtype='<U17')

```

1.0.7 Problem

Use these features in a LogisticRegression model. Create a barplot of the top five and bottom five coefficients of the model and their feature name. What do these mean?

1.0.8 n-Grams

```

1 vect = CountVectorizer(ngram_range=(2,2))

```

```

1 ex = train_slice.project_essay_1[:5]
2 y = train_slice.project_is_approved[:5]

```

```

1 vect.fit(ex)

```

```

CountVectorizer(analyzer='word', binary=False, decode_error='strict',
dtype=<class 'numpy.int64'>, encoding='utf-8', input='content
',
lowercase=True, max_df=1.0, max_features=None, min_df=1,
ngram_range=(2, 2), preprocessor=None, stop_words=None,
strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',
tokenizer=None, vocabulary=None)

```

```
1 vect.get_feature_names()[:10]
```

```
['00 in',  
'20 and',  
'25 new',  
'30 minutes',  
'580 students',  
'5th grade',  
'80 of',  
'88 5th',  
'8th grade',  
'92 of']
```

```
1 pipe = make_pipeline(TfidfVectorizer(min_df = 3),  
    ↳ LogisticRegression())  
2 params = {'logisticregression__C': [0.1, 1.0, 5.0, 10.0, 50, 100,  
    ↳ 1000], 'tfidfvectorizer__ngram_range': [(1,1), (1,2), (1,3)]}
```

```
1 grid = GridSearchCV(pipe, param_grid=params)
```

```
1 grid.fit(ex, y)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/model_selection/_split.  
py:605: Warning: The least populated class in y has only 2 members,  
which is too few. The minimum number of members in any class cannot be  
less than n_splits=3.  
% (min_groups, self.n_splits)), Warning)
```

```
GridSearchCV(cv=None, error_score='raise',  
            estimator=Pipeline(memory=None,  
            steps=[('tfidfvectorizer', TfidfVectorizer(analyzer='word',  
binary=False, decode_error='strict',  
dtype=<class 'numpy.int64'>, encoding='utf-8', input='content  
,  
lowercase=True, max_df=1.0, max_features=None, min_df=3,  
ngram_range=(1, 1), norm='l2', preprocessor=None, smooth_i...  
ty='l2', random_state=None, solver='liblinear', tol=0.0001,  
verbose=0, warm_start=False))]),  
fit_params=None, iid=True, n_jobs=1,  
param_grid={'logisticregression__C': [0.1, 1.0, 5.0, 10.0, 50,  
100, 1000], 'tfidfvectorizer__ngram_range': [(1, 1), (1, 2), (1, 3)]},  
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',  
scoring=None, verbose=0)
```

```
1 grid.best_score_
```

```
0.6
```

```
1 grid.best_params_
```

```
{'logisticregression__C': 0.1, 'tfidfvectorizer__ngram_range': (1, 1)}
```

```
1 scores = grid.cv_results_['mean_train_score'].reshape(-1,3).T
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/utils/deprecation.py
:122: FutureWarning: You are accessing a training score ('
mean_train_score'), which will not be available by default any more in
0.21. If you need training scores, please set return_train_score=True
warnings.warn(*warn_args, **warn_kwargs)
```

```
1 import seaborn as sns
```

```
1 plt.figure(figsize = (10, 5))
2 sns.heatmap(scores, annot=True, xticklabels=params['
    ↳ logisticregression__C'], yticklabels=params['
    ↳ tfidfvectorizer__ngram_range'])
3 plt.xlabel('C values')
4 plt.ylabel('n-gram values')
5 plt.title('Cross Validation Performance')
```

```
Text(0.5,1,'Cross Validation Performance')
```

